

SHUBHANG VAGVALA (NO SPONSORSHIP REQUIRED)

San Francisco Bay Area, California | (413) 362-0915 | shubhang680@gmail.com

LinkedIn: linkedin.com/in/sv68 | **GitHub:** github.com/shubhangv68 | **Portfolio:** shubhangv68.github.io/pf.github.io/

Summary

M.S. CS candidate (May 2025) with 3 years' industry experience in **backend and full-stack** engineering. Skilled in Java, Go, Spring Boot, React, Kafka, Redis, PostgreSQL, AWS, and Docker. Proven track record improving throughput, latency, and reliability on high-traffic systems; delivered **LLM-powered automation** and observability improvements.

Education

University of Massachusetts Amherst

Feb 2024 – May 2025

M.S. in Computer Science, GPA: 3.97

Coursework: Machine Learning, Artificial Intelligence, Distributed Operating Systems, Software Engineering

Experience

MediBuddy

Jun 2022 – Jan 2023

Software Engineer

Bengaluru, India

- Scaled **Spring Boot/Java microservices** across mobile and web; optimized DB indexing, connection pooling, and caching to improve API throughput by **60%** and reduce failures by **12%**.
- Built **asynchronous, Redis-backed APIs and workers** (queues, idempotency, retries, backpressure); cut server response times by **30%** and improved horizontal scalability.
- Shipped an **AI-powered FAQ assistant** with OpenAI and prompt-engineering guardrails; reduced support resolution time by **35%** and boosted customer satisfaction.
- Delivered **order-management features** with Spring Boot + React serving **500k+ users**; improved API latency from **2s → 0.3s** and reduced cancellations by **20%**.
- Integrated **Google Maps APIs** (geocoding, routing/ETA) with jitter handling and server-side caching; improved delivery accuracy by **15–20%**.

Swiggy

Jun 2021 – May 2022

Software Engineer

Bengaluru, India

- Built **real-time pricing microservices** in Go with Kafka and gRPC; applied circuit breakers, exponential backoff, and TTL caches to reduce latency by **25%** and improve pricing consistency by **20%**.
- Designed resilient **Kafka pipelines** (partitioning, consumer groups, schema versioning) with feature toggles and consumer-flag checks; maintained a **99.5% SLA** for business-critical services.
- Automated nightly **inventory audits** via Databricks jobs over CDC-replicated Hive tables; batched messages to Kafka with validation/reconciliation, eliminating **45%** of ETL inconsistencies.
- Built internal **Grafana dashboards** and observability (SLOs, traces, actionable alerts/runbooks); reduced on-call triage time by **10%**.

Projects

College Alumni Networking & Mentorship Platform

- Built a career guidance platform with the **MERN stack, TypeScript, Google APIs, SendGrid, and Redis**; implemented authentication, real-time scheduling, and notifications. Supported **600+ sessions**, reduced coordination overhead, and improved student access to mentors.

User-Item Interaction Recommender Pipeline

- Engineered a recommendation system with **distributed pipelines** processing **61M+ interactions**; implemented ingestion with **Spark**, feature transformations, and optimized batching for training/inference. Integrated collaborative filtering, content-based, and deep learning models for scalable personalization.

Technical Skills

Languages: Java, Go, Python, TypeScript, JavaScript (ES6+), SQL

Frameworks/Tech: Spring Boot, React, Node.js, gRPC, REST APIs, JWT/OAuth

GenAI/ML: OpenAI APIs, LangChain, RAG, FAISS/pgvector, PyTorch

Cloud/DevOps: AWS, Docker, Databricks, CI/CD (GitHub Actions, Jenkins), Vercel, Supabase, Git

Data Engineering: Kafka, Redis, PostgreSQL, MongoDB, Apache Hive, Spark, ETL Pipelines

Testing & Monitoring: Unit/Integration Testing, Postman, Grafana, SLOs/Tracing, TDD

Practices: Agile/Scrum, Microservices, System Design, Prompt Engineering