

SHUBHANG VAGVALA

San Francisco Bay Area, California | (413) 362-0915 | shubhangvagvala68@gmail.com

LinkedIn: linkedin.com/in/sv68 | GitHub: github.com/shubhangv68 | Portfolio:

shubhangv68.github.io/pf.github.io/

Summary

U.S. Citizen â Authorized to work in the U.S. without sponsorship.

Software engineer with experience building and operating backend services and data-intensive applications at scale.

Worked across production systems involving inventory, pricing, and search, with a focus on correctness, performance, and operational reliability. Comfortable designing backend workflows, asynchronous processing, and large-scale data handling, with exposure to performance evaluation and optimization in complex software systems.

Experience

MediBuddy | Software Engineer | Bengaluru, India | Jun 2022 â Dec 2023

- Improved address accuracy for fulfillment by building a geo-normalized, map-backed model using Java, Spring Boot,

Google Maps APIs, MySQL and Redis, increasing first-attempt success by 20% and reducing support tickets by 25%.

- Optimized Elasticsearch inverted-index queries to surface substitutes for unavailable items; instrumented discovery

events into an ML-ready pipeline with 99%+ T+1 availability, reducing search drop-offs by 15%.

- Strengthened platform reliability by enforcing timeout budgets, retry backoff and idempotency with async offloading

and caching; reduced peak-time API failures by 25% and cut production incidents by 15% across the core platform.

- Rebuilt the order management platform to unify Medicines delivery, Lab collection and in-center bookings by designing

a canonical order-state machine and backend aggregation with cursor-based pagination; enabled status-driven React

UI, cut p95 latency by 90%, improved vendor sync to 98%, and reduced cancellations and support load by 25%.

Swiggy | Software Engineer Intern | Bengaluru, India | Jun 2021 â May 2022

- Validated the Inventory Master Service after major changes by writing unit and integration tests for Kafka consumers,

gRPC endpoints, and HTTP APIs using an in-house Go framework; fixed issues and improved code coverage by 15%.

- Supported on-call engineers during year-end festive peak traffic by monitoring inventory and pricing systems, reviewing

dashboards and logs, and assisting with operational checks to ensure smooth handling of high order volumes.

- Built a scheduled nightly ETL cron job in Python to refresh inventory-driven pricing data across stores nationwide,

processing millions of records per run using CDC replicas, Databricks, Hive, and Kafka; improved pricing correctness

and reduced price-escalation incidents by 10% on a platform with 13M+ DAU.

- Built Grafana dashboards for services and pipelines I worked on, creating panels for API request volume and latency,

error rates, producer/consumer health, and batch job execution metrics to improve observability and support monitoring.

Education