

SHUBHANG VAGVALA

San Francisco Bay Area, California | (413) 362-0915 | shubhangvagvala68@gmail.com

LinkedIn: linkedin.com/in/sv68 | GitHub: github.com/shubhangv68 | Portfolio: shubhangv68.github.io/pf.github.io/

Summary

U.S. Citizen – Authorized to work in the U.S. without sponsorship.

Software engineer with experience building and operating backend services and data-intensive applications at scale. Worked across production systems involving inventory, pricing, and search, with a focus on correctness, performance, and operational reliability. Comfortable designing backend workflows, asynchronous processing, and large-scale data handling, with exposure to performance evaluation and optimization in complex software systems.

Experience

MediBuddy

Software Engineer

Jun 2022 – Dec 2023

Bengaluru, India

- Improved address accuracy for fulfillment by building a geo-normalized, map-backed model using **Java, Spring Boot, Google Maps APIs, MySQL and Redis**, increasing first-attempt success by **20%** and reducing support tickets by **25%**.
- Optimized **Elasticsearch** inverted-index queries to surface substitutes for unavailable items; instrumented discovery events into an ML-ready pipeline with **99%+ T+1 availability**, reducing search drop-offs by **15%**.
- Strengthened platform reliability by enforcing timeout budgets, retry backoff and idempotency with async offloading and caching; reduced peak-time API failures by **25%** and cut production incidents by **15%** across the core platform.
- Rebuilt the order management platform to unify Medicines delivery, Lab collection and in-center bookings by designing a **canonical order-state machine** and backend aggregation with cursor-based pagination; enabled **status-driven React UI**, cut p95 latency by **90%**, improved vendor sync to **98%**, and reduced cancellations and support load by **25%**.

Swiggy

Software Engineer Intern

Jun 2021 – May 2022

Bengaluru, India

- Validated the Inventory Master Service after major changes by writing unit and integration tests for **Kafka** consumers, **gRPC** endpoints, and **HTTP** APIs using an in-house **Go** framework; fixed issues and improved code coverage by **15%**.
- Supported on-call engineers during year-end festive peak traffic by monitoring inventory and pricing systems, reviewing dashboards and logs, and assisting with operational checks to ensure smooth handling of high order volumes.
- Built a scheduled nightly ETL cron job in **Python** to refresh inventory-driven pricing data across stores nationwide, processing **millions of records per run** using CDC replicas, **Databricks, Hive**, and **Kafka**; improved pricing correctness and reduced price-escalation incidents by **10%** on a platform with **13M+ DAU**.
- Built **Grafana** dashboards for services and pipelines I worked on, creating panels for API request volume and latency, error rates, producer/consumer health, and batch job execution metrics to improve observability and support monitoring.

Education

University of Massachusetts Amherst

Feb 2024 – May 2025

M.S. in Computer Science, GPA: 3.97

Coursework: Distributed Operating Systems, Software Engineering, Artificial Intelligence, Systems in Data Science

Birla Institute of Technology and Science Pilani

Aug 2018 – May 2022

B.E. in Electronics and Communication Engineering, GPA: 3.61

Coursework: Object-Oriented Programming, Data Structures and Algorithms, Microprocessors, Computational Logic

Technical Skills

Programming Languages: Java, Python, Go, TypeScript, JavaScript, SQL, Bash.

Backend Development: Spring Boot, REST APIs, gRPC, GraphQL, Node.js, Express.js.

Frontend Development: React, Next.js, Redux, Axios, HTML/CSS, Swagger/OpenAPI, JWT.

Databases and Caching: PostgreSQL, MySQL, MongoDB, Redis, DynamoDB.

Data Engineering: ETL pipelines, Apache Kafka, Apache Spark, Databricks, Apache Hive.

DevOps & Cloud: Docker, Git, GitHub Actions, Jenkins, Maven, Gradle, EC2, S3, Lambda, IAM, CloudWatch, CI/CD.

Testing and Automation: JUnit, Mockito, Jest, React Testing Library, Playwright.

Practices & Miscellaneous: Agile/Scrum, SDLC, Grafana, Kibana.

Projects

Optimizing Streaming Data Processing with Apache Flink

Designed and evaluated a **work-stealing scheduler** for **distributed stream processing** in **Apache Flink** to address load imbalance and data skew; processed **61M+ NYC For-Hire Vehicle events**, implemented dynamic task redistribution, and benchmarked against Flink's default scheduler with gains in runtime, throughput, and CPU efficiency.

College Alumni Networking and Mentorship Platform

Designed and built a backend-driven mentorship scheduling platform using **TypeScript** and **Node.js**, implementing **JWT**-based authentication and **Redis**-backed asynchronous job queues for session scheduling and notifications; deployed containerized services on **AWS**, supporting **600+** mentor-mentee sessions while reducing manual coordination overhead.